

Report: AdaAX - Explaining RNNs with Adaptive Automata (Paper Structure)

Abstract

Recurrent Neural Networks (RNNs) are effective for sequential data but lack transparency. This paper proposes **AdaAX**, a method to explain RNNs by constructing a Deterministic Finite Automaton (DFA). Unlike prior methods that fix state partitions early, AdaAX forms DFA states *adaptively*. It identifies fine-grained “core sets” based on RNN transition patterns and merges them strategically, allowing a trade-off between the explanation’s fidelity (accuracy) and complexity (size). Experiments demonstrate AdaAX achieves higher fidelity with significantly smaller DFAs compared to baselines.

1. Introduction

- **Problem:** RNNs are powerful but function as “black boxes,” making it hard to understand or trust their decision-making process. Interpretable models are needed.
- **Proposed Solution:** Use a DFA as an interpretable proxy model for an RNN. States in the DFA abstract RNN hidden states, and transitions follow input symbols. Paths (patterns) to accepting states explain predictions.
- **Limitations of Existing Work:** Current DFA extraction methods often pre-partition the RNN’s hidden state space, leading to DFAs that are either inaccurate (low fidelity) or too large and complex to understand.
- **Contribution (AdaAX):**
 - A novel DFA extraction method using **adaptive states** formed by merging fine-grained core sets.
 - Decouples pattern identification (high fidelity) from state formation (controlled complexity).
 - Provides a mechanism to explicitly trade fidelity for lower complexity.
 - Achieves superior performance (higher fidelity, lower complexity) experimentally.

2. Preliminaries

- **Recurrent Neural Network (RNN):** Processes sequences $x = (x_1, \dots, x_T)$, $x_t \in \Sigma$ (alphabet), computing hidden states $h_t = g(h_{t-1}, x_t)$ and a final output $y = f(h_T)$. \mathbb{H} denotes the hidden state space.
- **Deterministic Finite Automaton (DFA):** A tuple $\mathcal{H} = (Q, \Sigma, \delta, q_0, F)$:
 - Q : Finite set of states.
 - Σ : Alphabet (same as RNN).
 - δ : Transition function ($Q \times \Sigma \rightarrow Q$).
 - q_0 : Start state (representing RNN’s h_0).
 - $F \subseteq Q$: Set of accepting states.

- **RNN Explanation via DFA:** The DFA \mathcal{H} explains the RNN \mathcal{R} if its state transitions and acceptance behavior approximate the RNN’s hidden state dynamics and final predictions.
- **Patterns:** Input sequences p such that $\delta(q_0, p) \in F$. They represent inputs leading to the target prediction.
- **Problem Definition:** Given an RNN \mathcal{R} and data \mathcal{D} , learn a DFA \mathcal{H} that maximizes **fidelity** and minimizes **complexity** (size $|Q|$).
 - **Fidelity:** Measures prediction agreement between \mathcal{H} and \mathcal{R} .

$$fidelity(\mathcal{H}) = \frac{\sum_{x \in \mathcal{D}} \mathbb{I}(\mathcal{R}(x) = \mathcal{H}(x))}{|\mathcal{D}|}$$

(Eq. 2)

- **Accepting States (F):** Often correspond to RNN hidden states leading to a specific class prediction.

$$F_{RNN} = \{h \in \mathbb{H} \mid f(h, x) = 1, \forall x \in \Sigma\}$$

(Eq. 1)

F_{DFA} contains abstract states representing F_{RNN} .

3. Related Work

(This section summarizes the context inferred from the paper’s motivation)

* Existing methods for extracting DFAs from RNNs often rely on clustering RNN hidden states (e.g., using K-means) *before* learning transitions. * This pre-partitioning can be suboptimal, as clusters based solely on proximity might not align well with the actual transition dynamics learned by the RNN. * Such methods can result in low-fidelity DFAs or require a very large number of states (high complexity) to capture the RNN’s behavior accurately.

4. The AdaAX Method

AdaAX employs a three-step process with adaptive state formation:

4.1 Step 1: Clustering (Initial Grouping)

- Collect hidden states from the RNN using training data \mathcal{D} .
- Perform an initial, coarse clustering (e.g., K-means) on the hidden states \mathbb{H} .
- Treat the start state (h_0) and accepting states (F_{RNN}) as distinct initial groups.
- *Purpose:* Primarily for efficiency in the pattern extraction step, not for defining final DFA states.

4.2 Step 2: Pattern Extraction (Backward Search & Core Sets)

- Performs a backward, depth-first search from the accepting states F_{RNN} towards the start state h_0 .

- Identifies **Core Sets**: For a focal set of states C and an input symbol x , the core set consists of preceding states h such that $g(h, x) \in C$.
 - Concept of preceding states $P(C)$:

$$P(C) = \{h \in \mathbb{H} \mid g(h, x) \in C, x \in \Sigma\}$$

(Based on Eq. 3)

- Core sets group states based on *shared transition behavior*, providing finer granularity than initial clusters.
- Traces paths (sequences of symbols) back to h_0 , defining **patterns**.
- **Pruning**: Patterns with low support (frequency in \mathcal{D}) below a threshold θ can be removed.
 - Pattern Support:

$$supp_{\mathcal{D}}(p) = \frac{\sum_{x \in \mathcal{D}} y(p, x)}{|\mathcal{D}|}$$

(Definition 2.4)

4.3 Step 3: Consolidation (DFA Construction & Merging)

- Builds the DFA \mathcal{H} iteratively by adding extracted patterns (typically sorted by support).
- Incorporates core sets and transitions from each pattern into the DFA.
- **Adaptive State Merging**: To control complexity, newly added core sets (q_t) are evaluated for merging with existing DFA states ($S \in Q_t$).
 - Find **Neighboring States** $\mathcal{N}(q_t, Q_t)$:

$$\mathcal{N}(q_t, Q_t) = \{S \in Q_t \mid d(q_t.h, S.h) < \tau\}$$

(Eq. 5)

where $q_t.h$ is the RNN hidden state value corresponding to q_t (related to Eq. 4: $q.h = f(f(f(h_0, p_1), p_2) \dots, p_l)$) and τ is a distance threshold.

- Merge q_t with the closest neighbor $S \in \mathcal{N}(q_t, Q_t)$ *only if* the estimated drop in fidelity is below a user-defined threshold Δ .
- Merging combines prefixes and handles outgoing transitions intelligently.
- This merging process forms the final **adaptive states** of the DFA, balancing fidelity and complexity.

5. Experiments

- **Setup**: AdaAX compared against baseline DFA extraction methods. LSTMs trained on various datasets.
- **Datasets**: Included synthetic data (e.g., Tomita grammars, other regular languages) and real-world data (e.g., Yelp reviews, MIMIC-III health records, Educational Process Mining).

- **Results:**
 - **Fidelity vs. Complexity:** AdaAX consistently produced DFAs with higher fidelity for a given complexity (number of states) or significantly lower complexity for comparable fidelity, compared to baselines.
 - **Effectiveness of Merging:** The consolidation step effectively reduced DFA size while preserving high fidelity.
 - **Sensitivity Analysis:** AdaAX showed better sensitivity in identifying “flip points” – minimal input changes causing prediction flips.
 - **Case Studies:** Demonstrated utility in understanding model behavior on specific datasets (e.g., diagnosing RNN failures).

6. Conclusion

AdaAX presents a novel approach for extracting explanatory DFAs from RNNs. By introducing **adaptive states** formed through identifying fine-grained **core sets** and then consolidating them via a fidelity-controlled **merging** process, AdaAX overcomes limitations of prior methods. It generates more accurate (higher fidelity) and simpler (lower complexity) explanations, providing a valuable tool for interpreting RNN behavior.